

[◀ Back to Week 2](#)[☰ Lessons](#)[Prev](#)[Next](#)

Programming Assignment: SQL for Data Science Assignment

You have not submitted. You must earn 17/21 points to pass.

Deadline Pass this assignment by August 27, 11:59 PM PDT

Instructions

[My submission](#)

[Discussions](#)

Frontmatter

You will use the command line to complete this assignment.

If you are using the virtual machine, open a Terminal.

Some relevant linux commands and notes for this assignment:

- `cd directory_name`: change directory to `directory_name`
- `ls`: list files in the current directory
- `sqlite3 reuters.db < sql_filename`: execute the queries in the file `sql_filename` against the database `reuters.db`. This command assumes `reuters.db` and `sql_filename` are in your current directory.
- Pressing the tab key will attempt to autocomplete any

How to submit

When you're ready to submit, you can upload files for each part of the assignment on the "My submission" tab.

command or filename you are currently typing.

- For Problem 2, you will use matrix.db instead of reuters.db (Some students do not notice the switch.)

As always, before you start the assignment, get the most recent course materials by issuing the following command: **git pull**

For most of these problems, you will use the reuters.db database consisting of a single table:

```
1 frequency(docid, term, count)
```

where docid is a document identifier corresponding to a particular file of text, term is an English word, and count is the number of the occurrences of the term within the document indicated by docid.

Many questions ask you to count the number of records returned by a query. Perhaps the easiest way to count the number of records returned by a query Q is to write Q as a subquery:

```
1 SELECT count(*) FROM (  
2     SELECT ...  
3 ) x;  
4
```

(In SQLite, the alias "x" is not required, but in other dialects of SQL it is. So we've included it here.)

Problem 1: Inspecting the Reuters Dataset and Basic Relational Algebra

(a) select: Write a query that is equivalent to the following relational algebra expression.

$\sigma_{\text{docid}=10398_txt_earn}(\text{frequency})$

What to turn in: Run your query against your local database and determine the number of records returned. Save that value in a file *part_a.txt* and upload the file as your answer.

(b) select project: Write a SQL statement that is equivalent to the following relational algebra expression.

$\pi_{\text{term}}(\sigma_{\text{docid}=10398_txt_earn \text{ and } \text{count}=1}(\text{frequency}))$

What to turn in: Run your query against your local database and determine the number of records returned. Save that value in a file *part_b.txt* and upload the file as your answer.

(c) union: Write a SQL statement that is equivalent to the following relational algebra expression. (Hint: you can use the UNION keyword in SQL)

$\pi_{\text{term}}(\sigma_{\text{docid}=10398_txt_earn \text{ and } \text{count}=1}(\text{frequency})) \cup$
 $\pi_{\text{term}}(\sigma_{\text{docid}=925_txt_trade \text{ and } \text{count}=1}(\text{frequency}))$

What to turn in: Run your query against your local database and determine the number of records returned. Save that value in a file *part_c.txt* and upload the file as your answer.

(d) count: Write a SQL statement to count the number of unique documents containing the word "law" or containing the word "legal" (If a document contains both law and legal, it should only be counted once)

What to turn in: Run your query against your local database and determine the number of records returned. Save that value in a file *part_d.txt* and upload the file as your answer.

(e) big documents Write a SQL statement to find all documents that have more than 300 total terms, including duplicate terms. (Hint: You can use the HAVING clause, or you can use a nested query. Another hint: Remember that the *count* column contains the term frequencies, and you want to consider duplicates.)
(docid, term_count)

What to turn in: Run your query against your local database and determine the number of records returned. Save that value in a file part_e.txt and upload the file as your answer.

(f) two words: Write a SQL statement to count the number of unique documents that contain both the word 'transactions' and the word 'world'. (Hint: Find the docs that contain one word and the docs that contain the other word separately, then find the intersection.)

What to turn in: Run your query against your local database and determine the number of records returned as described above. Save that value in a file part_f.txt and upload the file as your answer.

Problem 2: Matrix Multiplication in SQL

Recall from lecture that a sparse matrix has many positions with a value of zero.

Systems designed to efficiently support sparse matrices look a lot like databases: They represent each cell as a record (i,j,value).

The benefit is that you only need one record for every non-zero element of a matrix.

For example, the matrix

--	--	--

0	2	-1
1	0	0
0	0	-3
0	0	0

can be represented as a table

row #	column #	value
0	1	2
0	2	-1
1	0	1
2	2	-3

Take a minute to make sure you understand how to convert back and forth between these two representations.

Now, since you can represent a sparse matrix as a table, it's reasonable to consider whether you can express matrix multiplication as a SQL query and whether it makes sense to do so.

Within matrix.db, there are two matrices A and B represented as follows:

```

1  A(row_num, col_num, value)
2
3  B(row_num, col_num, value)
4
```

The matrix A and matrix B are both square matrices with 5 rows and 5 columns each.

(g) multiply: Express $A \times B$ as a SQL query, referring to the class lecture for hints.

What to turn in: Save the value of cell (2,3) in a file part_g.txt and upload the file as your answer. The file should only contain the value of the cell at position (2,3).

If you're wondering why matrix multiply in a database might be a good idea, consider that advanced databases execute queries in parallel automatically. So it can be quite efficient to process a very large sparse matrix --- millions of rows or columns --- in a database. But a word of warning: In a job interview, don't necessarily tell them you recommend implementing linear algebra in a database. You won't be wrong, but the interviewer might not understand databases as well as you now do, and therefore won't understand when and why this is a good idea. Just say you have done some experiments using databases for analytics.

Problem 3: Working with a Term-Document Matrix

The reuters dataset can be considered a *term-document matrix*, which is an important representation for text analytics.

The reuters dataset can be considered a *term-document matrix*, which is an important representation for text analytics.

Each row of the matrix is a *document vector*, with one column for every term in the entire corpus. Naturally, some documents may not contain a given term, so this matrix is rather sparse. The value in each cell of the matrix is the term frequency. (You'd often want this value to be a *weighted* term frequency, typically using "tf-idf": *term frequency - inverse document frequency*. But we'll stick with the raw frequency for now.)

What can you do with the term-document matrix D ? One thing you can do is compute the similarity of documents. Just multiply the matrix with its own transpose $S = DD^T$, and you have an (unnormalized) measure of similarity.

The result is a square document-document matrix, where each cell represents the similarity. Here, similarity is pretty simple: if two documents both contain a term, then the score goes up by the product of the two term frequencies. This score is equivalent to the dot product of the two document vectors.

To normalize this score to the range 0-1 and to account for relative term frequencies, the *cosine similarity* is perhaps more useful. The cosine similarity is a measure of the angle between the two document vectors, normalized by magnitude. You just divide the dot product by the magnitude of the two vectors. However, we would need a power function (x^2 , $x^{(1/2)}$) to compute the magnitude, and sqlite has built-in support for only very basic mathematical functions. It is not hard to extend sqlite to add functions that you need, but we won't be doing that in this assignment.

(h) similarity matrix: Write a query to compute the similarity matrix DD^T . (Hint: The transpose is trivial -- just join on columns to columns instead of columns to rows.) The query could take some time to run if you compute the entire result. But notice that you don't need to compute the similarity of both (doc1, doc2) and (doc2, doc1) -- they are the same, since similarity is symmetric. If you wish, you can avoid this wasted work by adding a condition of the form $a.docid < b.docid$ to your query. (But the query still won't return immediately if you try to compute every result -- don't expect otherwise.)

What to turn in: Create a file `part_h.txt` containing the similarity value of the two documents '10080_txt_crude' and '17035_txt_earn'. Upload this file as your answer.

You can also use this similarity metric to implement some primitive search capabilities. Consider a keyword query that you might type into Google: It's a bag of words, just like a document (typically a keyword query will have far fewer terms than a document, but that's ok).

So if we can compute the similarity of two documents, we can compute the similarity of a query with a document. You can imagine taking the union of the keywords represented as a small set of (docid, term, count) tuples with the set of all documents in the corpus, then recomputing the similarity matrix and returning the top 10 highest scoring documents.

(i) keyword search: Find the best matching document to the keyword query "washington taxes treasury". You can add this set of keywords to the document corpus with a union of scalar queries:

```
1 SELECT * FROM frequency
2 UNION
3 SELECT 'q' as docid, 'washington' as term, 1 as count
4 UNION
5 SELECT 'q' as docid, 'taxes' as term, 1 as count
6 UNION
7 SELECT 'q' as docid, 'treasury' as term, 1 as count
8
```

Then, compute the similarity matrix again, but filter for only similarities involving the "query document": docid = 'q'. Consider creating a view of this new corpus to simplify things.

What to turn in: Create a file part_i.txt containing the maximum similarity score between the keyword query among all documents. Your SQL query should return a list of (docid, similarity) pairs, but you will submit only include a single number: the highest similarity score in the list.

